

```
#Importing the libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
sns.set()
import warnings
warnings.filterwarnings('ignore')
```

```
#loading the dataset
df = pd.read_csv('usedCars.csv')
df.head()
```

	Id	Company	Model	Variant
FuelType \				
0	555675	MARUTI SUZUKI	CELERIO(2017-2019)	1.0 ZXI AMT 0
PETROL				
1	556383	MARUTI SUZUKI	ALTO	LXI
PETROL				
2	556422	HYUNDAI	GRAND I10	1.2 KAPPA ASTA
PETROL				
3	556771	TATA	NEXON	XT PLUS
PETROL				
4	559619	FORD	FIGO EXI DURATORQ	1.4 DIESEL

	Colour	Kilometer	BodyStyle	TransmissionType	ManufactureDate
ModelYear \					
0	Silver	33197	HATCHBACK	NaN	2018-02-01
2018					
1	Red	10322	HATCHBACK	Manual	2021-03-01
2021					
2	Grey	37889	HATCHBACK	Manual	2015-03-01
2015					
3	A Blue	13106	HATCHBACK	NaN	2020-08-01
2020					
4	Silver	104614	HATCHBACK	Manual	2010-11-01
2010					

	CngKit	Price	Owner	DealerState
DealerName \				
0	NaN	5.75 Lakhs	1st Owner	Karnataka
Cars				Top Gear
1	NaN	4.35 Lakhs	1st Owner	Karnataka
Renew 4 u Automobiles PVT Ltd				
2	NaN	4.7 Lakhs	1st Owner	Karnataka
Anant Cars Auto Pvt Ltd				
3	NaN	9.9 Lakhs	1st Owner	Karnataka
Adeep Motors				
4	NaN	2.7 Lakhs	2nd Owner	Karnataka
Zippy				

Automart

	City	Warranty	QualityScore
0	Bangalore	1	7.8
1	Bangalore	1	8.3
2	Bangalore	1	7.9
3	Bangalore	1	8.1
4	Bangalore	0	7.5

## Some Numerical Information about the Data

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1064 entries, 0 to 1063
```

```
Data columns (total 19 columns):
```

#	Column	Non-Null	Count	Dtype
0	Id	1064	non-null	int64
1	Company	1064	non-null	object
2	Model	1064	non-null	object
3	Variant	1064	non-null	object
4	FuelType	1063	non-null	object
5	Colour	1064	non-null	object
6	Kilometer	1064	non-null	int64
7	BodyStyle	1064	non-null	object
8	TransmissionType	350	non-null	object
9	ManufactureDate	1064	non-null	object
10	ModelYear	1064	non-null	int64
11	CngKit	22	non-null	object
12	Price	1064	non-null	object
13	Owner	1064	non-null	object
14	DealerState	1064	non-null	object
15	DealerName	1064	non-null	object
16	City	1064	non-null	object
17	Warranty	1064	non-null	int64
18	QualityScore	1064	non-null	float64

```
dtypes: float64(1), int64(4), object(14)
```

```
memory usage: 158.1+ KB
```

```
df.nunique()
```

Id	1064
Company	23
Model	218
Variant	575
FuelType	5
Colour	76
Kilometer	1006
BodyStyle	10

TransmissionType	9
ManufactureDate	162
ModelYear	17
CngKit	2
Price	367
Owner	4
DealerState	10
DealerName	57
City	11
Warranty	2
QualityScore	43

dtype: int64

```
df.isnull().sum()
```

Id	0
Company	0
Model	0
Variant	0
FuelType	1
Colour	0
Kilometer	0
BodyStyle	0
TransmissionType	714
ManufactureDate	0
ModelYear	0
CngKit	1042
Price	0
Owner	0
DealerState	0
DealerName	0
City	0
Warranty	0
QualityScore	0

dtype: int64

## Data Preprocessing

```
# Define a function for convert price to float value
def convert_amount(amount_str):
    if "Lakhs" in amount_str:
        return float(amount_str.replace(' Lakhs', '').replace(',', '')) * 100000
    else:
        return float(amount_str.replace(',', ''))

df['Price'] = df['Price'].apply(convert_amount)

# Drop Unnecessary Columns
df = df.drop(['Id', 'CngKit', 'TransmissionType'], axis=1)
```

```

# Apply fillna on FuelType column
df['FuelType'] = df['FuelType'].fillna('PETROL')

# Define a function for reduce unique values of categorical columns
def unique_reduce(x, dic):
    if x in dic.keys():
        return x
    else :
        return 'Others'

fuel_dic = dict(df['FuelType'].value_counts().head(2))
df['FuelType'] = df['FuelType'].apply(lambda x : unique_reduce(x,
fuel_dic))
df['FuelType'].value_counts()

FuelType
PETROL      671
DIESEL      365
Others       28
Name: count, dtype: int64

company_dic = df['Company'].value_counts().head(7)
df['Company'] = df['Company'].apply(lambda x : unique_reduce(x,
company_dic))
df['Company'].value_counts()

Company
MARUTI SUZUKI    252
Others           219
HYUNDAI          199
HONDA            126
MAHINDRA          96
TATA              60
FORD              58
TOYOTA           54
Name: count, dtype: int64

color_dic = df['Colour'].value_counts().head(5)
df['Colour'] = df['Colour'].apply(lambda x : unique_reduce(x,
color_dic))
df['Colour'].value_counts()

Colour
Others    341
White    289
Silver   134
Grey     127
Red      109
Black     64
Name: count, dtype: int64

```

```
body_dic = df['BodyStyle'].value_counts().head(3)
df['BodyStyle'] = df['BodyStyle'].apply(lambda x : unique_reduce(x,
body_dic))
df['BodyStyle'].value_counts()
```

```
BodyStyle
HATCHBACK    423
SUV           304
SEDAN         262
Others        75
Name: count, dtype: int64
```

```
dealer_dic = df['DealerState'].value_counts().head(4)
df['DealerState'] = df['DealerState'].apply(lambda x :
unique_reduce(x, dealer_dic))
df['DealerState'].value_counts()
```

```
DealerState
Others        355
Delhi         196
Maharashtra   194
Karnataka     165
Haryana       154
Name: count, dtype: int64
```

```
city_dic = df['City'].value_counts().head(6)
df['City'] = df['City'].apply(lambda x : unique_reduce(x, city_dic))
df['City'].value_counts()
```

```
City
Others        234
Delhi         196
Bangalore     165
Gurgaon       154
Pune          147
Noida         95
Kolkata       73
Name: count, dtype: int64
```

```
def dealer_category(x):
    if x in ['Car Estate', 'Carz Villa', 'SUSHIL CARS PVT. LTD',
'Taneja Fourwheels', 'Car Choice Exclusif', 'Fast Wheels Cars', 'K.S.
Motors']:
        return 'Dealer Type1'
    elif x in ['Sai Motors', 'LUXMI CARS GURGAON', 'Guru Kripa
Motors', 'Star Auto India']:
        return 'Dealer Type2'
    elif x in ['Shree Radha Krishna Motors', 'Shiv Auto Wings',
'Mahindra First Choice Wheels Ltd', 'Noida Car Point ll', 'Jeen Mata
Motors', 'OM Motors', 'Sireesh Auto Pvt Ltd', 'Max Motors', 'Adeep
Motors']:
```

```

        return 'Dealer Type3'
    elif x in ['Ikka Motors', 'Pitbox Motors', 'Noida Car Ghar',
               'Cardiction', 'Car&Bike Superstore Pune', 'Zippy Automart', 'SK
               Associates', 'Anant Cars Auto Pvt Ltd', 'DrivUS Motorcorp', 'Vinayak
               Autolink Private Limited'] :
        return 'Dealer Type4'
    elif x in ['Prestige Autoworld Pvt Ltd', 'PROPEL MOTORS', 'Royal
               Motors (Prop. Auto Carriage Pvt Ltd)', 'MM Motors', 'Sri Vaishnavi
               Cars', 'Top Gear Cars', 'Renew 4 u Automobiles PVT Ltd'] :
        return 'Dealer Type5'
    else :
        return 'Other Dealers'

df['DealerName'] = df['DealerName'].apply(lambda x :
dealer_category(x))
df['DealerName'].value_counts()

DealerName
Other Dealers      209
Dealer Type4       199
Dealer Type5       188
Dealer Type3       182
Dealer Type1       180
Dealer Type2       106
Name: count, dtype: int64

# change type of owner column to int and remove 3 and 4 values
df['Owner'] = df['Owner'].str[0]
df['Owner'] = df['Owner'].map(int)
df= df[df['Owner'] < 3 ]

```

## Data Visualization

```

# Define list of Continuous columns Names
continuous = ['Price', 'Kilometer', 'QualityScore']

# Distribution of Categorical Features
def plot_continious_distribution(df, column):

    width_ratios = [2, 4]
    gridspec_kw = {'width_ratios':width_ratios}
    fig, ax = plt.subplots(1, 2, figsize=(12, 6), gridspec_kw =
gridspec_kw)
    fig.suptitle(f' {column} ', fontsize=20)

    sns.boxplot(df[column], ax=ax[0])
    ax[0].set_title('Boxplot Chart')
    ax[0].set_ylabel(column)

    sns.histplot(x = df[column], kde=True, ax=ax[1], multiple =

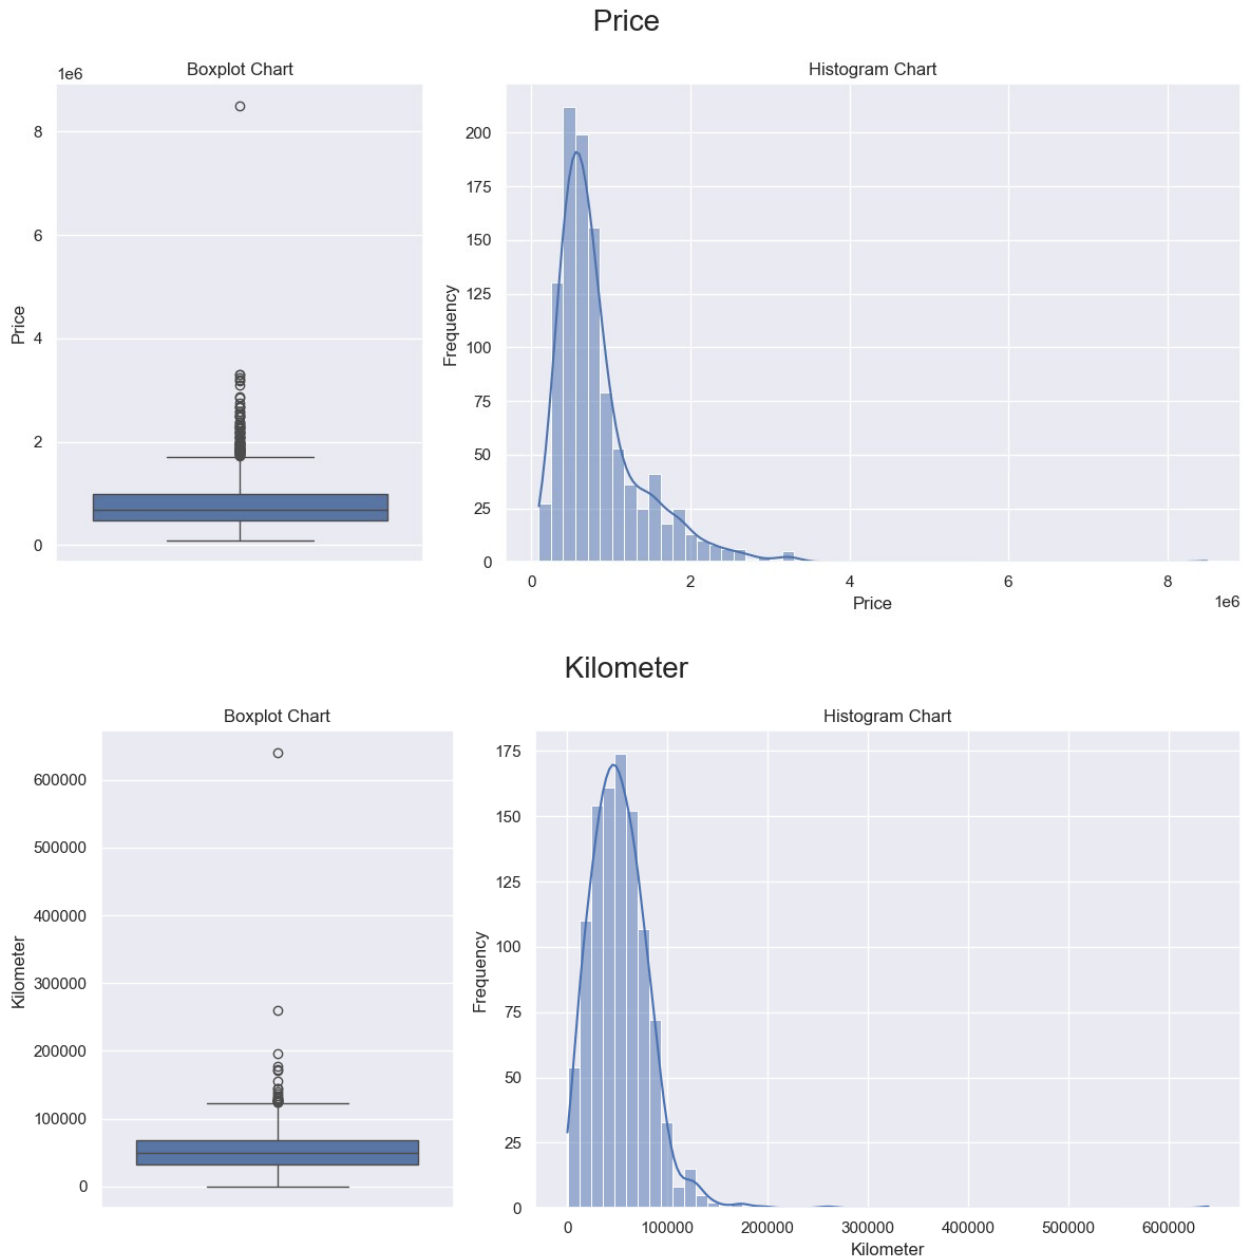
```

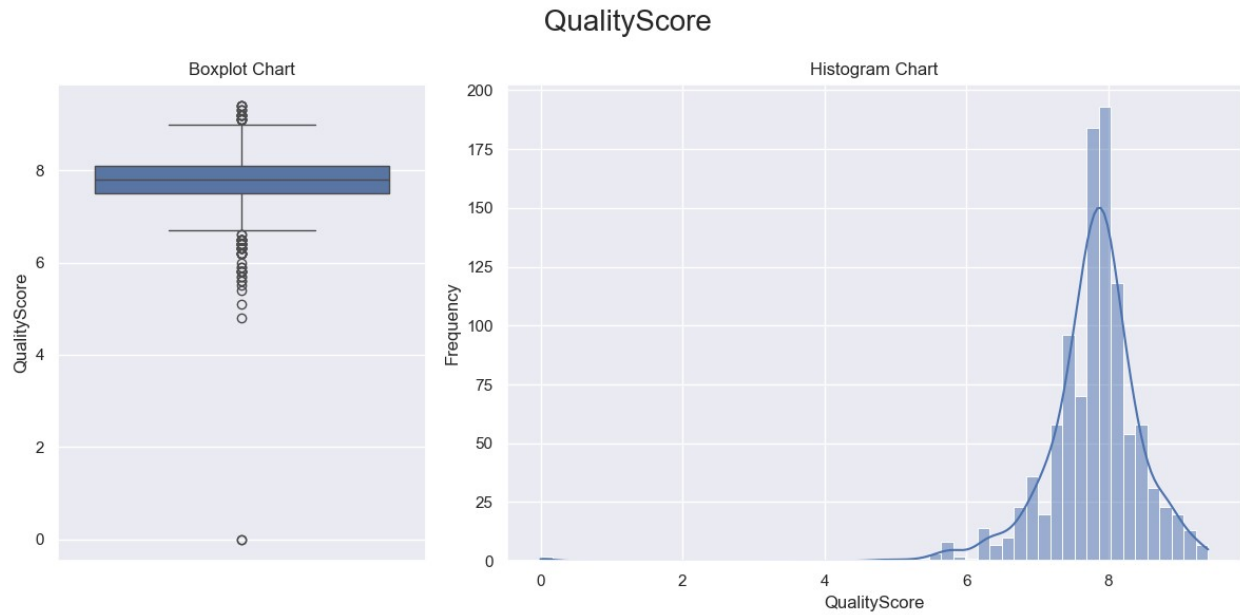
```

'stack', bins=55)
ax[1].set_title('Histogram Chart')
ax[1].set_ylabel('Frequency')
ax[1].set_xlabel(column)

plt.tight_layout()
plt.show()
for conti in continuous :
    plot_continious_distribution(df, conti)

```





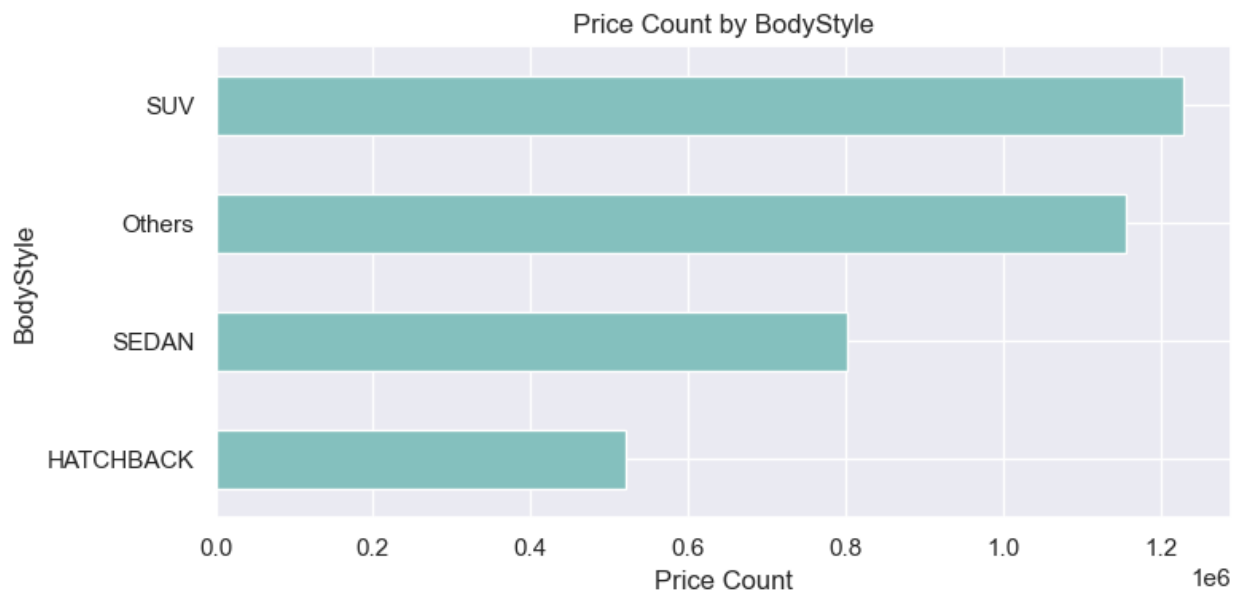
*# Define a Function for Barh Plot*

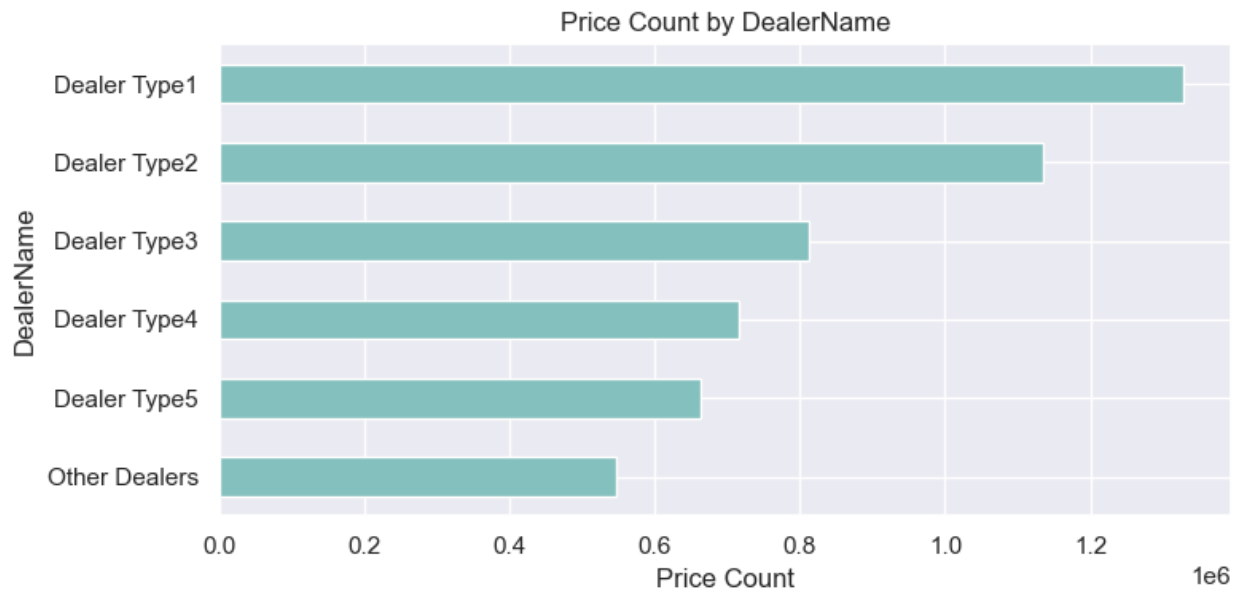
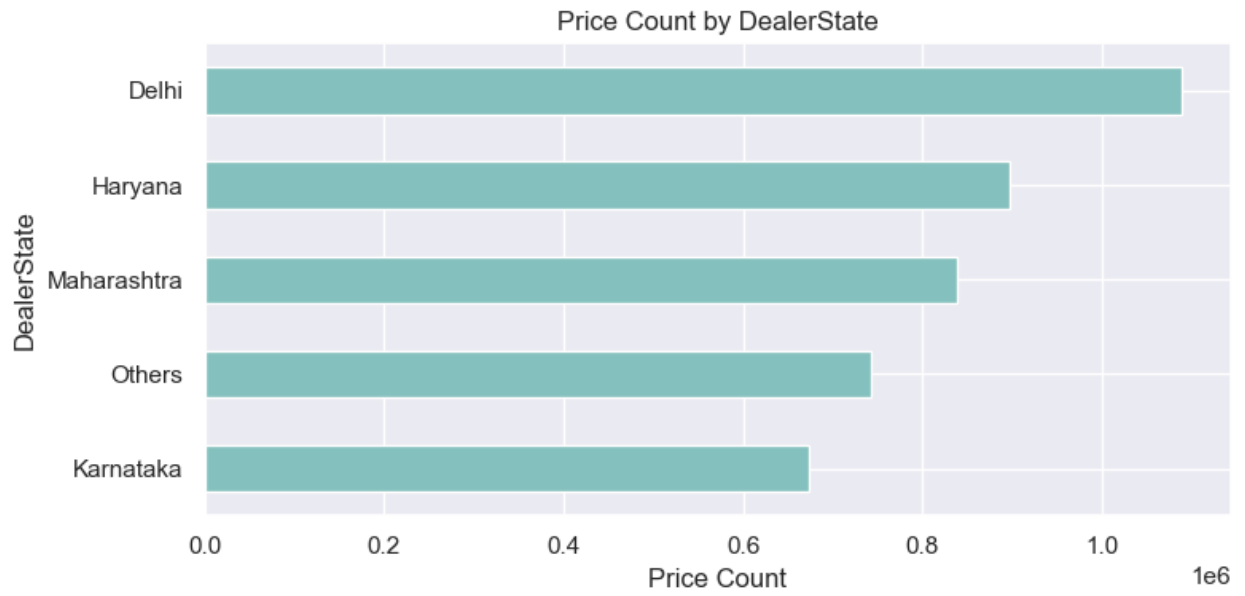
```
def bar_plot(x, y, df):
    barh = df.groupby([x])[y].mean()
    barh.sort_values(ascending=True, inplace=True)
    barh.plot(kind='barh', color = '#84c0be', figsize=(8,4))
    plt.title(f'{y} Count by {x}')
    plt.xlabel(f'{y} Count')
```

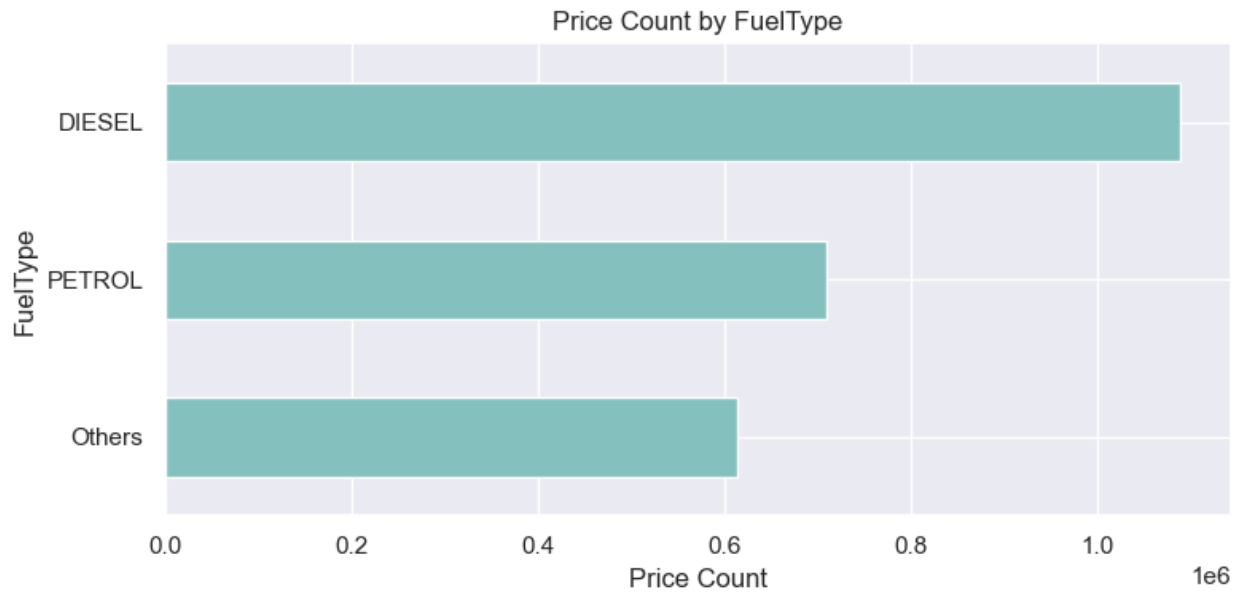
```
    plt.tight_layout()
    plt.show()
```

```
bar_plot('Company', 'Price', df)
bar_plot('BodyStyle', 'Price', df)
bar_plot('DealerState', 'Price', df)
bar_plot('DealerName', 'Price', df)
bar_plot('FuelType', 'Price', df)
bar_plot('City', 'Price', df)
```









## Data Preprocessing

```
from sklearn.preprocessing import LabelEncoder, StandardScaler

# Initialize StandardScaler
stc = StandardScaler()
# Initialize LabelEncoder
le = LabelEncoder()

stc_cols = ['Kilometer', 'Price', 'QualityScore']
dum_cols = ['Company', 'FuelType', 'Colour', 'BodyStyle',
            'DealerState', 'DealerName', 'City']
```

```

# Apply Standard Scaler to the selected columns
df[stc_cols] = stc.fit_transform(df[stc_cols])
# Apply Get Dummies to the selected columns
df = pd.get_dummies(df, columns=dum_cols)

```

## Training and Evaluating Different Models

```

from sklearn.model_selection import train_test_split

x = df.drop(['Price', 'Model', 'Variant', 'ManufactureDate'], axis=1)
y = df['Price']

x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, random_state=42)

#Importing the Libraries
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import LogisticRegression
from sklearn.datasets import make_classification
from sklearn.model_selection import GridSearchCV
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import VotingRegressor
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error, r2_score

# List of Mdels to Try
models = [
    ('Linear Regression', LinearRegression()),
    ('Ridge Regression', Ridge()),
    ('Decision Tree', DecisionTreeRegressor()),
    ('Random Forest', RandomForestRegressor()),
    ('Gradient Boosting', GradientBoostingRegressor()),
    ('K-Nearest Neighbors', KNeighborsRegressor()),
    ('XGB Regressor', XGBRegressor())
]

# Train and evaluate each model
for name, model in models:
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    print(f'{name}: Mean Squared Error = {round(mse,3)}, R-squared =
{round(r2, 3)}')

Linear Regression: Mean Squared Error = 0.243, R-squared = 0.66
Ridge Regression: Mean Squared Error = 0.243, R-squared = 0.66

```

Decision Tree: Mean Squared Error = 0.51, R-squared = 0.286  
Random Forest: Mean Squared Error = 0.184, R-squared = 0.742  
Gradient Boosting: Mean Squared Error = 0.209, R-squared = 0.708  
K-Nearest Neighbors: Mean Squared Error = 0.275, R-squared = 0.615  
XGB Regressor: Mean Squared Error = 0.212, R-squared = 0.703

```
from sklearn.model_selection import GridSearchCV

# Define the parameter grid to search
param_grid = {
    'n_estimators': [160, 150, 170],
    'max_depth': [None, 10, 20],
}

# Initialize the Random Forest Regressor
rf_model_tuned = RandomForestRegressor(random_state=42)

# Initialize GridSearchCV
grid_search = GridSearchCV(rf_model_tuned, param_grid, cv=3,
    scoring='neg_mean_squared_error', n_jobs=-1, verbose=True)

# Fit the grid search to the data
grid_search.fit(x_train, y_train)

# Get the best parameters
rf_best_params = grid_search.best_params_

# Retrain the model with the best parameters
rf_model_best = RandomForestRegressor(**rf_best_params,
    random_state=42)
rf_model_best.fit(x_train, y_train)

# Predict using the updated features
y_pred_best = rf_model_best.predict(x_test)

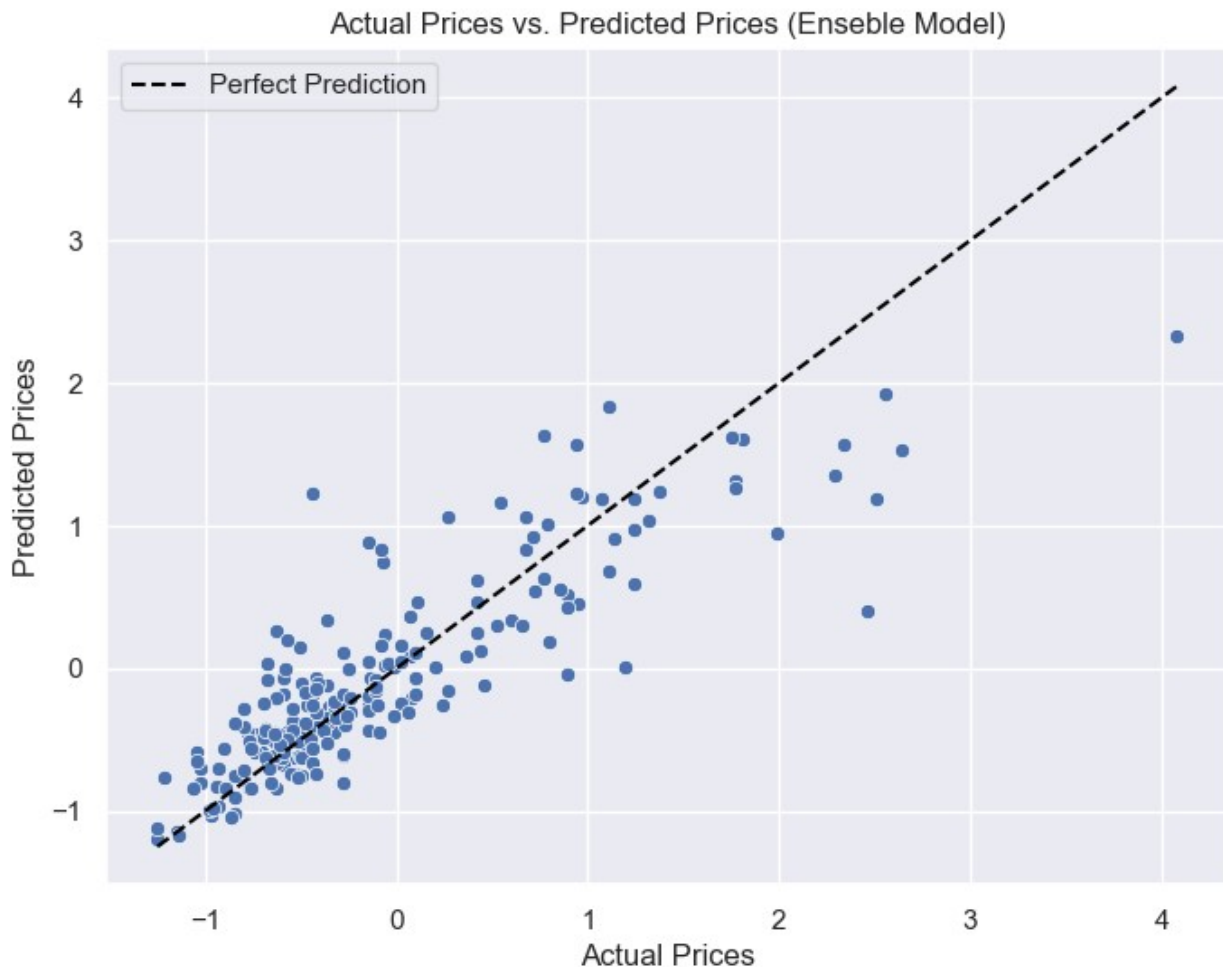
Fitting 3 folds for each of 9 candidates, totalling 27 fits

# Evaluate the tuned Random Forest model
mse_best = mean_squared_error(y_test, y_pred_best)
r2_best = r2_score(y_test, y_pred_best)

print(f'Best Parameters: {rf_best_params}')
print(f'Mean Squared Error (Tuned Random Forest): {round(mse_best, 3)}')
print(f'R-squared (Tuned Random Forest): {round(r2_best, 3)}')

Best Parameters: {'max_depth': 10, 'n_estimators': 150}
Mean Squared Error (Tuned Random Forest): 0.18
R-squared (Tuned Random Forest): 0.748
```

```
# Visualize the Predicted Prices Against the Actual Prices
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test, y=y_pred_best)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
linestyle='--', color='black', label='Perfect Prediction')
plt.title('Actual Prices vs. Predicted Prices (Enseble Model)')
plt.ylabel('Predicted Prices')
plt.xlabel('Actual Prices')
plt.legend()
plt.show()
```



## Summary and Conclusion

In this project, I focused on predicting the prices of used cars in India using various data preprocessing techniques and a machine learning model. The steps and methodologies employed are as follows:

1. Data Cleaning and Preprocessing:
  - String to Numeric Conversion: Converted values in the price column from strings to numeric format.

- Handling Missing Values: Removed columns with a high number of null values to ensure data quality and consistency.
  - Categorical Simplification: Reduced the number of unique values in some categorical columns to simplify the model.
2. Data Visualization:
    - Created appropriate visualizations to explore and understand the data patterns and relationships, providing valuable insights into the dataset.
  3. Data Standardization and Label Encoding:
    - Performed data standardization to normalize the features.
    - Applied label encoding to convert categorical variables into numerical format.
  4. Model Training and Optimization:
    - Trained a Random Forest model on the processed dataset.
    - Optimized the model using Grid Search to improve accuracy.
  5. Model Evaluation:
    - The final model achieved an accuracy of 74.8%. Given the complexity of the model and the variability of the data, this result is considered satisfactory.

These steps ensured a comprehensive analysis and model training process, leading to a reasonably accurate prediction model for used car prices in Indi

*Developed by Hosein Mohammadi*

GitHub : <https://github.com/Hosein541>

Kaggle : <https://www.kaggle.com/hoseinnnnnnn>

Gmail : [Huseinmohammadi83@gmail.com](mailto:Huseinmohammadi83@gmail.com)